

系統程式設計

施吉昇
臺灣大學資訊工程系



Agenda

- Properties of Files and Directories
- File systems and Directory Structures
- Special access permission
- Ownership of new files

Properties of Files and Directories

Files in POSIX and Unix Systems

- Resources in operating systems can be classified into
 - computable resources: process, memory, etc.
 - non-computable resources: storage systems, hardware devices, communication links, etc.
- File is a general concept to mange non-computable resources.
 - Storage systems:
 - regular files, directories, etc.
 - Devices: keyboard, pointer devices, USB, etc.
 - /proc/*: interfaces to communicate with kernel
 - /etc/sockets: interface to communicate with processes (on different machines.)

File Types

- Regular Files: text, binary, etc.
- Directory Files: Only Kernel can update these files – { (filename, pointer) }.
- Character Special Files, e.g., tty, audio, etc.
- Block Special Files, e.g., disks, etc.
- FIFO – named pipes
- Sockets – not POSIX.1 or SVR4
 - SVR4 uses library of socket functions, instead.
 - 4.3+BSD has a file type of socket.
- Symbolic Links – not POSIX.1 or SVR4

Link

- Files are identified by i-nodes in the file systems.
 - It is a long string, usually represented by a long address, to be read by the operating system.
 - I-node is not designed for human users.
- Directories and filename are widely used by librarian/human users to organize large volumes of documents, long before there are computers.
 - Link is the pointer to bridge legible file name/ directory name to i-node.
 - One i-node can have more than one name, each of which is represented by one link.


Links: hard and soft

/usr/joe

foo

/usr/sue

bar



File i-node:
Reference = 2
(File i-node)

(Data structure to list the files in a directory)

- (Hard) Link

- Each directory entry creates a hard link of a filename to the i-node that describes the file's contents.

Symbolic Links (soft link)

(directory)

foo



a


(file)

testdir

- Initially introduced by 4.2BSD
- Provide a pointer to a file/directory, similar to the concept of alias.
 - It has a dedicated file type: symbolic link
 - It is represented as a regular file in the system.
 - The content of the file is the path of the targeted file/directory.
- The existence of the targeted file/directory is NOT validated by the file system and the link may
 - form a loop or
 - point to an invalid path for file/directory.


Links: hard and soft

/usr/joe



- Hard Link

- Each directory entry creates a hard link of a filename to the i-node that describes the file's contents.



- Symbolic Link (Soft Link)

- It is implemented as a file that contains a pathname.
- Filesize = pathname length
- Example: Shortcut on Windows

* Problem – infinite loop in tracing a path name with symbolic links – 4.3BSD, no 8 passings of soft links

* Dangling pointers

Hard Link vs. Soft Link

- Hard Link

- Cannot link to different mounted file system.
- is a different name for the same set of data blocks.
- Only superuser can create a link to a directory.

- Soft Link (or Symbolic link)

- It is implemented as a file that contains a pathname.
- File size = pathname length
- Can be a directory or file
- Is a pointer to a set of data blocks.
- File type is S_IFLINK

- What's the difference on their i-nodes?

File Information

- Three major functions:

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
int stat(const char *pathname, struct stat *buf);
int fstat(int filedes, struct stat *buf);
int lstat(const char *pathname, struct stat *buf);
int fstatat(int fd, const char *restrict pathname, struct
stat *restrict buf, int flag);
```

Differences on stat(), fstat(), lstat():

lstat() returns info regarding the symbolic link itself, instead of the referenced file, if it happens.

fstatat(): returns info for a pathname relative to an open directory represented by *fd* argument.

File Properties

- Differences on stat(), fstat(), lstat():

- lstat() returns info regarding the symbolic link, instead of the referenced file, if it happens.

```
struct stat {
    mode_t      st_mode; /* type & mode */
    ino_t       st_ino; /* i-node number */
    dev_t       st_dev; /* device no (file system) */
    dev_t       st_rdev; /* device no for special file */
    nlink_t     st_nlink; /* # of links */
    uid_t       st_uid; gid_t st_gid;
    off_t       st_size; /* sizes in bytes */
    struct timespec st_atime; /* last access time */
    struct timespec st_mtime; /* last modification time */
    struct timespec st_ctime; /* time for last status change */
    blksize_t     st_blk_size; /* best I/O block size */
    blkcnt_t      st_blocks; /* number of disk blocks allocate
};
```

File Types

- Program 4.1 – Page 96

- lstat() and Figure 4.1

```
<sys/stat.h>
#define S_IFMT      0xF000 /* type of file */
#define S_IFDIR     0x4000 /* directory */
#define S_ISDIR(mode) (((mode)&0xF000) == 0x4000)
#define S_ISDIR(mode) (((mode) & S_IFMT) == S_IFDIR)
```

- st_mode

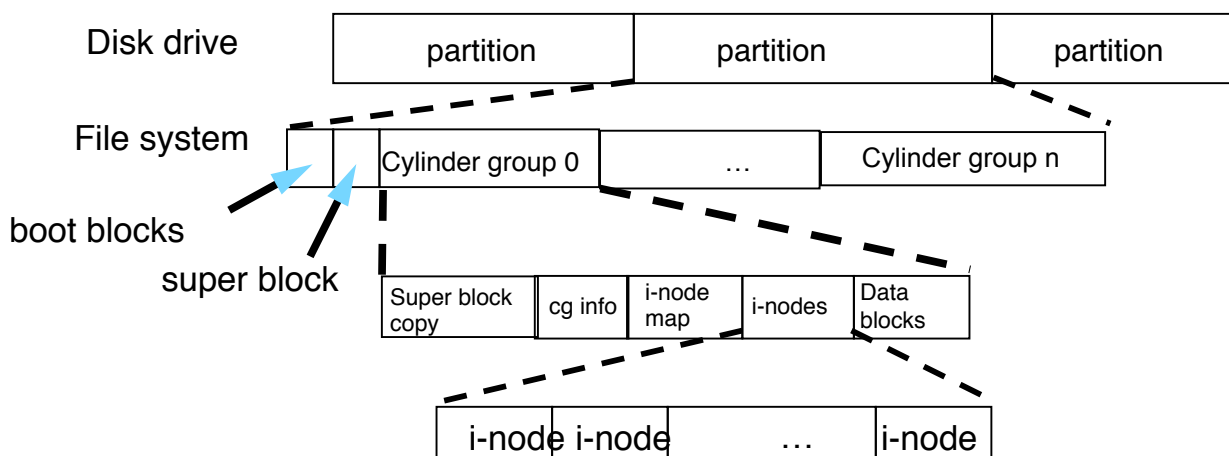
- Percentage of Files in a Medium-Sized System – Figure 4.4

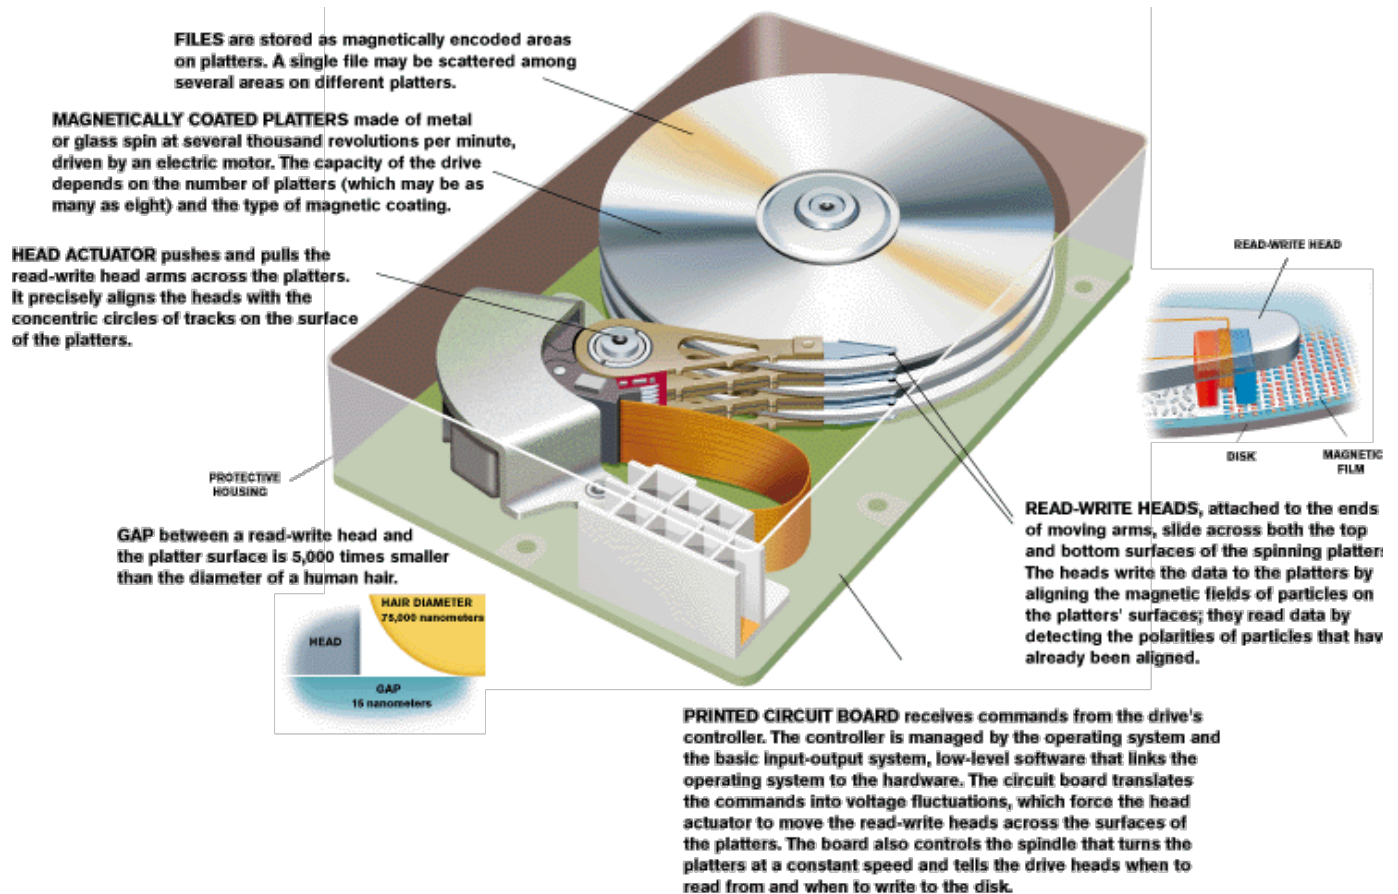
File Type	Count	Percentage
regular file	226,856	88.22%
directory	23,017	8.95%
symbolic link	6,442	2.51%
character special	447	0.17%
block special	312	0.12%
socket	69	0.03%
FIFO	1	0.00%

File systems and Directory Structure


File Systems Arrangement

- A hierarchical arrangement of directories and files – starting in root /
- Several Types, e.g., SVR4: Unix System V File Systems (S5), Unified File System (UFS) – Figure 2.6 (Page 39)
- System V:





Cylinder, track, cluster, and sector



Partition vs. Physical Devices

Disk /dev/hda: 123.5 GB, 123522416640 bytes
255 heads, 63 sectors/track, 15017 cylinders
Units = cylinders of 16065 * 512 = 8225280 bytes

Device	Boot	Start	End	Blocks	Id	System
/dev/hda1	*	1	13	104391	83	Linux
/dev/hda2		14	14895	119539665	83	Linux
/dev/hda3		14896	15017	979965	82	Linux swap

```
[root@oris /]# more /etc/mtab
/dev/hda1      /boot      ext3      rw 0 0
/dev/hda2      /          ext3      rw 0 0
none          /proc      proc      rw 0 0
none          /dev/pts   devpts    rw,gid=5,mode=620 0 0
usbdevfs      /proc/bus/usb usbdevfs  rw 0 0
none          /dev/shm   tmpfs     rw 0 0
/dev/sda1      /mnt/usbhd ext3      rw 0 0
```

Superblock in File System

- Each file system has different format and information:
 - Type: ext2, ext3, FAT, FAT32 etc.
 - Size: 5 GB, 10 GB, etc and
 - Status such as mount status.
- A superblock contains information about file system such as: File system type, Size, Status, and Information about other metadata structures.
- When superblock of a filesystem is lost, we are in big trouble. So, you will find redundant superblocks in the system.
- Example: Linux with root permission

```
dumpe2fs /dev/sda1 | grep -i superblock
```

```
[root@oris sbin]# ./dumpe2fs /dev/sda1 | grep -i superblock
```

```
dumpe2fs 1.40.4 (31-Dec-2007)
```

```
Primary superblock at 0, Group descriptors at 1-7
```

```
Backup superblock at 32768, Group descriptors at 32769-32775
```

```
Backup superblock at 98304, Group descriptors at 98305-98311
```

```
Backup superblock at 163840, Group descriptors at 163841-163847
```

```
Backup superblock at 229376, Group descriptors at 229377-229383
```

```
Backup superblock at 294912, Group descriptors at 294913-294919
```

```
Backup superblock at 819200, Group descriptors at 819201-819207
```

```
Backup superblock at 884736, Group descriptors at 884737-884743
```

```
Backup superblock at 1605632, Group descriptors at 1605633-1605639
```

```
Backup superblock at 2654208, Group descriptors at 2654209-2654215
```

```
Backup superblock at 4096000, Group descriptors at 4096001-4096007
```

```
Backup superblock at 7962624, Group descriptors at 7962625-7962631
```

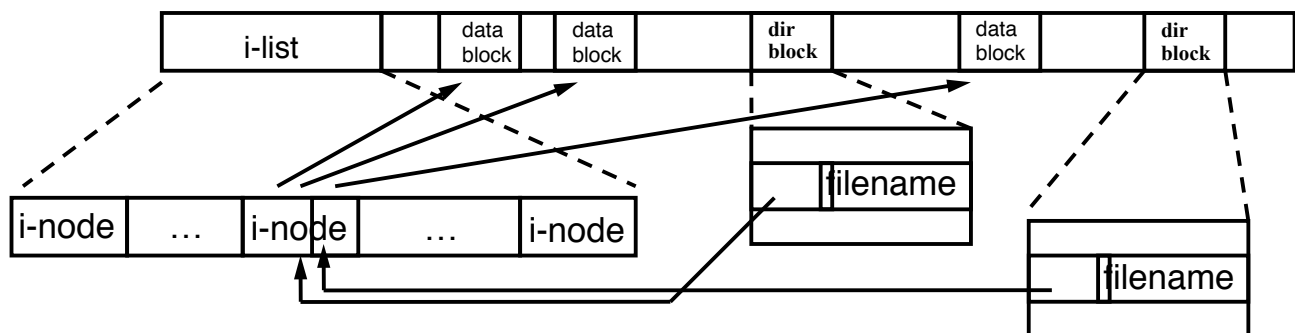
```
Backup superblock at 11239424, Group descriptors at 11239425-11239431
```

```
Backup superblock at 20480000, Group descriptors at 20480001-20480007
```

```
Backup superblock at 23887872, Group descriptors at 23887873-23887879
```

```
[root@oris sbin]#
```

i-node and data blocks



● i-node:

- Version 7: 64B, 4.3+BSD:128B, S5:64B, UFS:128B
- File type, access permission, file size, data blocks, etc.

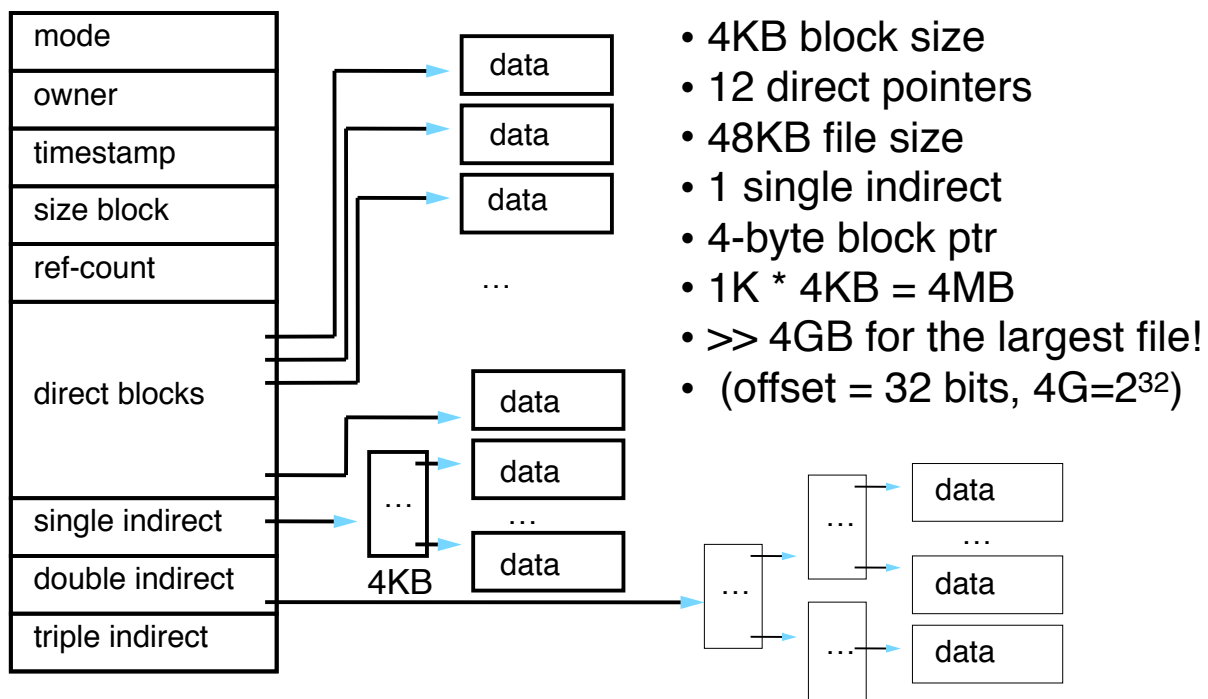
● Link count – hard links

- st_nlink in stat, LINK_MAX in POSIX.1
- Unlink/link a file

Moving files among directories

- When files are moved within the same mounted file system, no need to physically move the file.
 - Only directory block is updated.
- Otherwise, the files have to be physically moved.

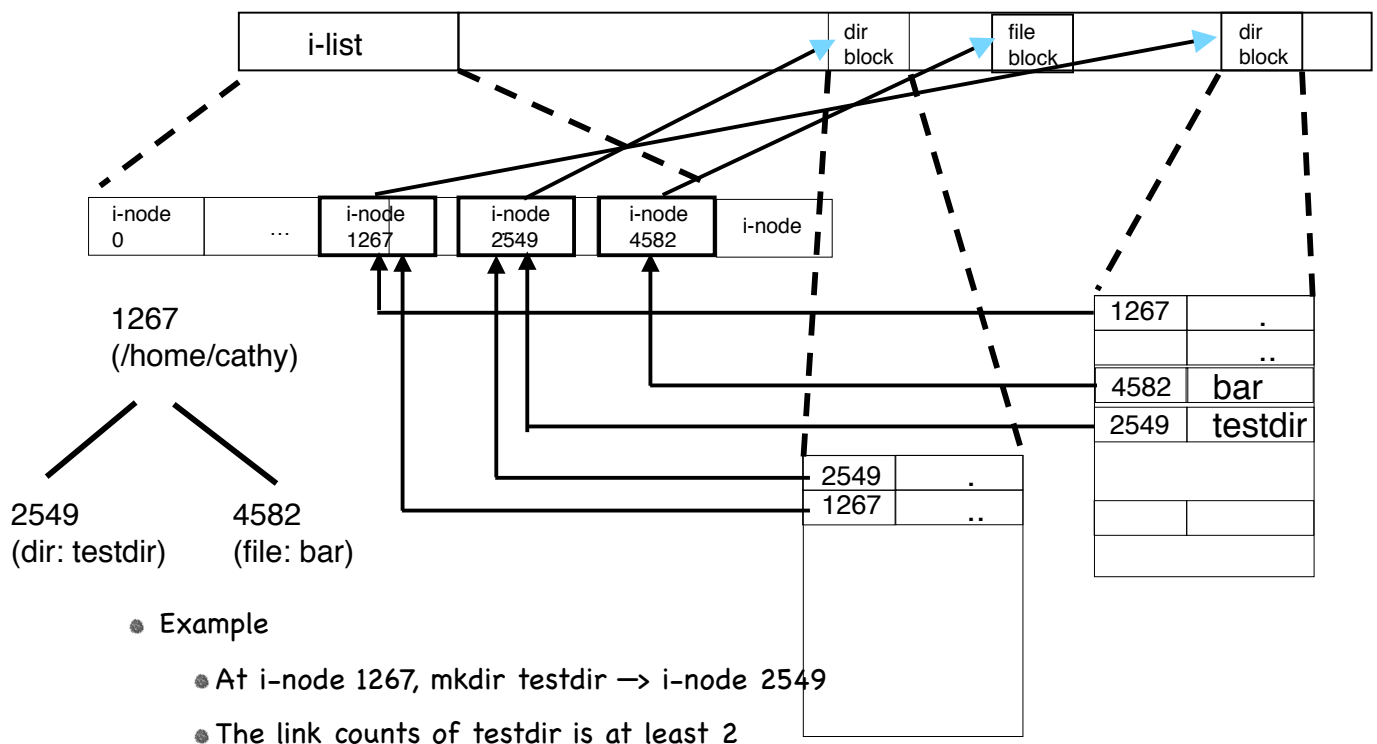
File System - 4.4BSD i-node



Link

- Files are identified by i-nodes in the file systems.
 - It is a long string, usually represented by a long address, to be read by the operating system.
 - I-node is not designed for human users.
- Directories and filename are widely used by librarian/human users to organize large volumes of documents, long before there are computers.
 - Link is the pointer to bridge legible file name/directory name to i-node.
 - One i-node can have more than one name, each of which is represented by one link.

Links to organize files/directories



- Command mv only modify dir entries!
- No directory entry points at any i-node residing at other file systems.

Functions – link and unlink

```
#include <unistd.h>
```

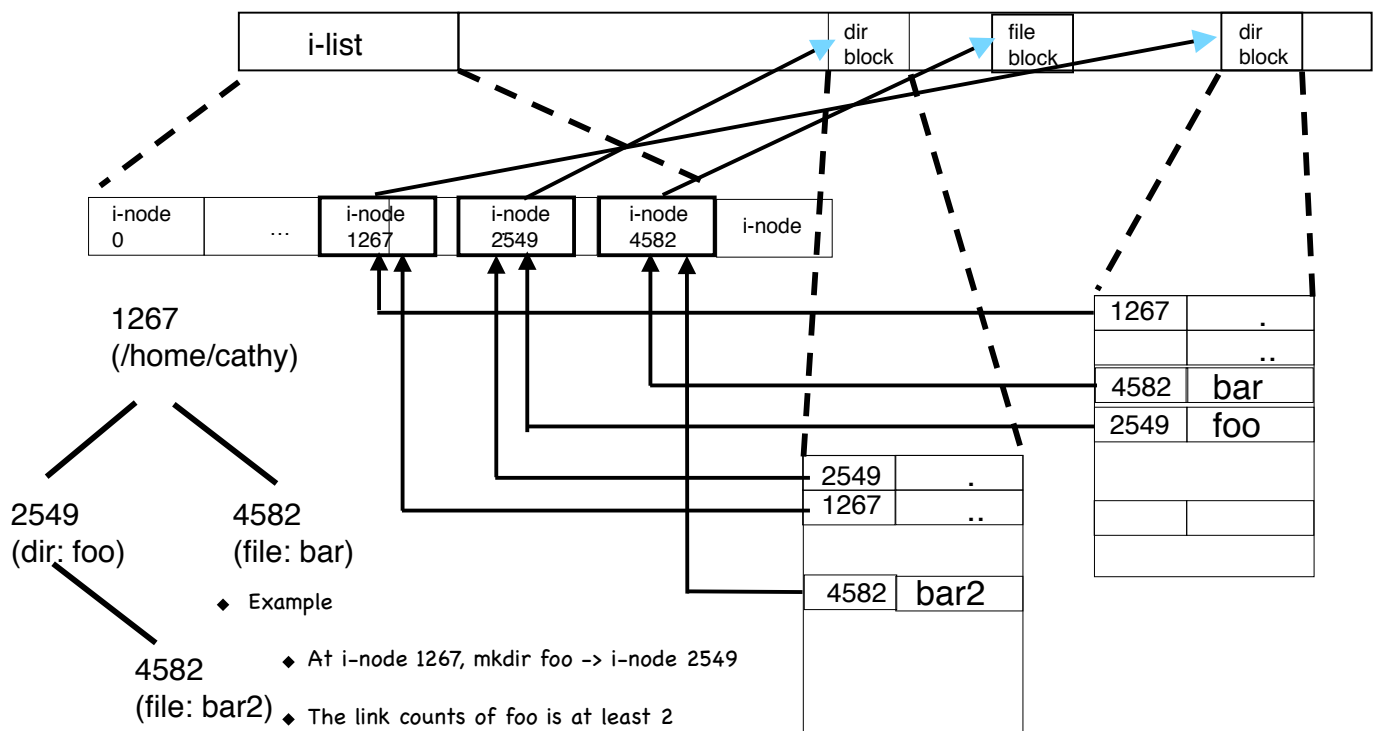
```
int link(const char *existingpath, const char *newpath);
```

```
int linkat(int efd, const char *existingpath, int nfd, const  
char *newpath, int flag);
```

```
int unlink(const char *pathname);
```

- Atomic action for link – hard link
- POSIX.1 allows linking across file systems (most systems do not support this implementation.)
- Only superusers could create a link to a dir
- Error if newpath exists
- Unlink – WX permission at the residing dir
- Remove the dir entry & delete the file if link count reaches zero and no one still opens the file (Remark: sticky bit & dir rights).

Links to organize files/directories



- ♦ Command mv only modify dir entries!
- ♦ No directory entry points at any i-node residing at other file systems.

Example for link and unlink

- Program 4.16 – Page 118
 - Open & unlink a file
- Unlink a file
 - When sticky bits set for a residing dir, we must have the write permission for the directory and either owner of the file, the dir, or super users.
 - If pathname is a symbolic link, it removes the symbolic link and does not change the count value of the referred i-node.

Functions – rename and remove

```
#include <stdio.h>
```

```
int remove(const char *pathname);
```

```
int rename(const char *oldname, const char *newname);
```

- remove =
 - unlink if pathname is a file.
 - rmdir if pathname is a dir. (ANSI C)
- Rename – ANSI C
 - File: both files, newname is removed first, WX permission for both residing directories
 - Directory: both dir, newname must be empty, newname could not contain oldname. E.g., rename("bar", "bar/foo").

Symbolic Links

- Goal

- Get around the limitations of hard links: (a) file system boundary (b) link to a dir.


- Initially introduced by 4.2BSD

- Example – ftw: file tree walk

- In `-s ../foo testdir`

- Figure 4.17 – functions follow slinks

- No functions which take filedes
- Example: `unlink(testdir)`



Files follow symbolic link

	Does not	Follows
access		Y
chdir		Y
chmod		Y
chown	Y	Y
creat		Y
exec		Y
lchown	Y	
link		Y
truncate		Y

	Does not	Follows
lstat	Y	
open		Y
opendir		Y
pathconf		Y
readlink	Y	
remove	Y	
rename	Y	
stat		Y
unlink	Y	

Create/Read Symbolic Links

```
#include <unistd.h>

int symlink(const char *actualpath, const char *sympath);

int readlink(const char *pathname, char *buf, int
    bufsize);
```

- actualpath does not need to exist!
 - They do not need to be in the same file system.
- readlink is an action consisting of open, read, and close – not null terminated.

Hard Link/Soft Link in Windows-based OS

- NTFS Hard link, similar to hard link in POSIX
 - Available since Windows 2000
 - Link to files on the same volumes
 - fsutil hardlink create <NewFileName> <ExistingFileName>
- NTFS Junction,
 - Can link files and directories to different volumes on the same host.
 - linkd mydesktop user profile/desktop
 - linkd mydesktop /d
- Symbolic link, similar to soft link in POSIX
 - Available since NTFS 5.0, i.e., Windows Vista
 - Can also point to a file or remote SMB network path
 - mklink [[/D] | [/H] | [/J]] link target
 - /D – Creates a directory symbolic link. Default is a file symbolic link.
 - /H – Creates a hard link instead of a symbolic link.
 - /J – Creates a Directory Junction.

Discussion

- Alice creates a symbolic link and removes the target file
 - `ln -s bar foo`
 - `rm bar`
- and, in her program, she writes
 - `open("foo", O_RDONLY);`
- What will `open()` return?

Functions – `mkdir` and `rmdir`

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
int mkdir(const char *pathname, mode_t mode);
```

- `umask`, UID/GID setup (Sec 4.6)
- From 4.2BSD & SVR3 (SVR4 – inheritance of the `S_ISGID` bit)
- At least one of the execute bits are enabled.

```
#include <unistd.h>
```

```
int rmdir(const char *pathname);
```

- An empty dir is deleted.
- Link count reaches zero, and no one still opens the dir.

Functions – opendir, readdir, rewinddir, closedir

```
#include <sys/types.h>
#include <dirent.h>
DIR *opendir(const char *pathname);
struct dirent *readdir(DIR *dp);
void rewinddir(DIR *dp);
int closedir(DIR *dp);
```


- Only the kernel can write to a dir!!!
- WX for creating/deleting a file!
- Not POSIX standard but XSI extension for UNIX

Functions – opendir, readdir, rewinddir, closedir

- dirent struct is implementation-dependent, e.g.,

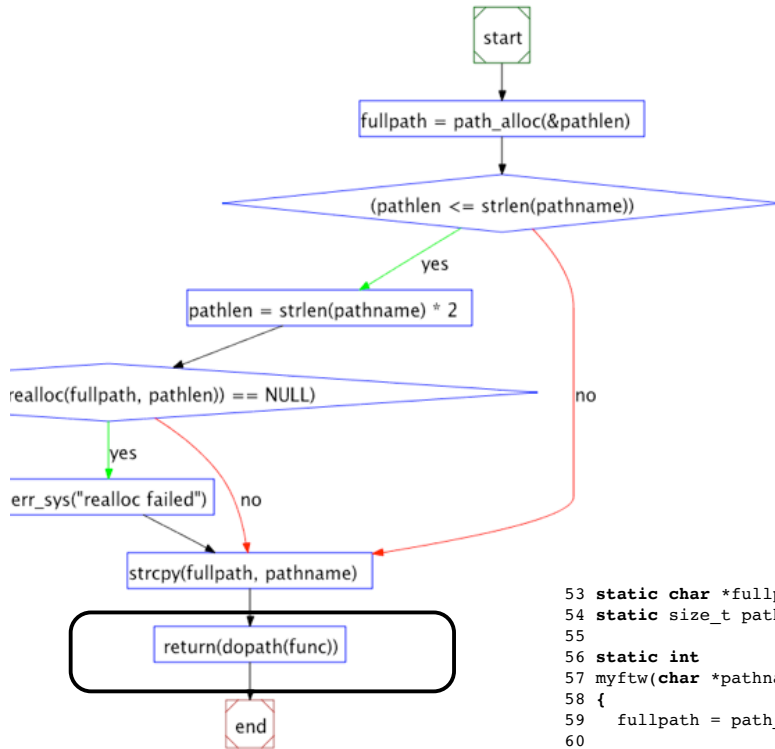
```
struct dirent {
    ino_t d_ino;  /* not in POSIX.1 */
    char d_name[NAME_MAX+1];
} /* fpathconf() */
```

- Program 4.22 – Pages 132–133
 - ftw/nftw – recursively traversing the file system



	Da	Db	Fd	Fe	Ff	Dc	Fg
nreg	0	0	1	2	2	2	3
ndir	1	2	2	2	2	3	3






```

53 static char *fullpath; /* contains full pathname for every file */
54 static size_t pathlen;
55
56 static int /* we return whatever func() returns */
57 myftw(char *pathname, Myfunc *func)
58 {
59     fullpath = path_alloc(&pathlen); /* malloc PATH_MAX+1 bytes */
60     /* ({Prog pathalloc}) */
61     if (pathlen <= strlen(pathname)) {
62         pathlen = strlen(pathname) * 2;
63         if ((fullpath = realloc(fullpath, pathlen)) == NULL)
64             err_sys("realloc failed");
65     }
66     strcpy(fullpath, pathname);
67     return(dopath(func));
  
```



```

76 static int /* we return whatever func() returns */
77 dopath(Myfunc* func)
78 {
79     struct stat statbuf;
80     struct dirent *dirp;
81     DIR *dp;
82     int ret, n;
83
84     if (lstat(fullpath, &statbuf) < 0) /* stat error */
85         return(func(fullpath, &statbuf, FTW_NS));
86     if (S_ISDIR(statbuf.st_mode) == 0) /* not a directory */
87         return(func(fullpath, &statbuf, FTW_F));
88
89     /* It's a directory. First call func() for the directory,
90      * then process each filename in the directory.
91      */
92     if ((ret = func(fullpath, &statbuf, FTW_D)) != 0)
93         return(ret);
94
95     n = strlen(fullpath);
96     if (n + NAME_MAX + 2 > pathlen) { /* expand path buffer */
97         pathlen *= 2;
98         if ((fullpath = realloc(fullpath, pathlen)) == NULL)
99             err_sys("realloc failed");
100     }
101     fullpath[n++] = '/';
102     fullpath[n] = 0;
103
104     if ((dp = opendir(fullpath)) == NULL) /* can't read directory */
105         return(func(fullpath, &statbuf, FTW_DNR));
106
107     while ((dirp = readdir(dp)) != NULL) {
108         if (strcmp(dirp->d_name, ".") == 0 ||
109             strcmp(dirp->d_name, "..") == 0)
110             continue; /* ignore dot and dot-dot */
111         strcpy(&fullpath[n], dirp->d_name); /* append name after "/" */
112         if ((ret = dopath(func)) != 0) /* recursive */
113             break; /* time to leave */
114     }
115     fullpath[n-1] = 0; /* erase everything from slash onward */
116
117     if (closedir(dp) < 0)
118         err_ret("can't close directory %s", fullpath);
119     return(ret);
120 }
121
  
```



	Da	Db	Fd	Fe	Ff	Dc	Fg
nreg	0	0	1	2	2	2	3
ndir	1	2	2	2	2	3	3

Functions – chdir, fchdir, getcwd

```
#include <unistd.h>
```

```
int chdir(const char *pathname);
```

```
int fchdir(int filedes);
```

- fchdir – not POSIX.1
- chdir must be built into shells!
- The kernel only maintains the i-node number and dev ID for the current working directory!
- Per-process attribute – working dir!
- Program 4.23 – Page 135
 - chdir, which does perform like chdir in command line.
 - How can you write a program to change the directory of your shell environment?

Functions - getcwd

```
#include <unistd.h>
```

```
char *getcwd(char *buf, size_t size);
```

- The system keeps the current working directory in the form of i-node number. getcwd() return the full path.
- The buffer must be large enough, or an error returns!
- chdir follows symbolic links, and getcwd has no idea of symbolic links!

• Program 4.24 - Page 137

- getcwd

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[])
{
    char save_cwd[1024];
    char buf[1024];

    /* Save the current working directory */
    if (getcwd(save_cwd, sizeof(save_cwd)) == NULL)
        perror("getcwd error");
    printf("before chdir, cwd = %s\n", save_cwd);

    /* change to a different working directory */
    if (chdir("/") < 0)
        perror("chdir error");

    if (getcwd(buf, sizeof(buf)) == NULL)
        perror("getcwd error");
    printf("after chdir, cwd = %s\n", buf);

    /* Return to earlier working directory */
    if (chdir(save_cwd) < 0)
        perror("chdir error");

    if (getcwd(buf, sizeof(buf)) == NULL)
        perror("getcwd error");
    printf("previous chdir, cwd = %s\n", buf);
    exit(0);
}
```

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[])
{
    char buf[1024];
    int fd;

    /* Save the current working directory */
    if ((fd = open(".", O_RDONLY)) < 0)
        perror("open error");

    if (getcwd(buf, sizeof(buf)) == NULL)
        perror("getcwd error");
    printf("before chdir, cwd = %s\n", buf);

    /* change to a different working directory */
    if (chdir("/") < 0)
        perror("chdir error");

    if (getcwd(buf, sizeof(buf)) == NULL)
        perror("getcwd error");
    printf("after chdir, cwd = %s\n", buf);

    /* Return to earlier working directory */
    if (fchdir(fd) < 0)
        perror("fchdir error");

    if (getcwd(buf, sizeof(buf)) == NULL)
        perror("getcwd error");
    printf("previous chdir, cwd = %s\n", buf);
    exit(0);
}
```

File Times

```
#include <sys/types.h>
#include <utime.h>
int utime(const char *pathname, const struct utimbuf *times);
int utimes(const char *pathname, const struct timeval
times[2]);
```

- time values are in seconds since the Epoch

- times = null -> set as the current time

- Effective UID = file UID or W right to the file

```
struct utimbuf {
    time_t actime;
    time_t modtime;
}
```

- times != null -> set as requested

- (Effective UID = file UID or superuser) and W right to the file.

- Program 4.6 (zap.c) - Page 118, utime

```
struct timeval {
    time_t      tv_sec;    /* seconds */
    suseconds_t tv_usec;   /* microseconds */
};
```

- What about ctime?

File Times

- Three Time Fields:

Field	Description	Example	ls option
st_atim	last-access-time	Read	-u
st_mti	last-modification-	Write	default
st_ctim	last-i-node-change-	chmod, chown	-c

- Figure 4.20 - Effect of functions on times

- Changing the access permissions, user ID, link count, etc, only affects the i-node!

- ctime is modified automatically! (stat, access)

- Changed by the following functions: chmod(2), chown(2), link(2), mknod(), pipe(), unlink(2), utime(), and write().

On Windows platforms, st_ctime refers to the creation time of the file

- Example Question: reading/writing a file only affects the file, instead of the residing dir (Fig4.20).

Discussion

- When you change the name of a file, which time attributes are changed?
- When you are asked to remove the 'core' or unused files which have not been read for one month, which time attributes should you use?

Access Permission

Access Permissions

- File Access Permission:
 - Permissions include read, write, execute, and others.
 - Permissions are granted according to user identity and user groups identity.

```
  8 7 6   5 4 3   2 1 0
|-----|-----|-----|
| user   group  other |
| r w x | r w x | r w x |
|-----|-----|-----|
```

Octal	Binary	Permission
0	000	no permission
1	001	execute
2	010	write
3	011	write and execute
4	100	read
5	101	read and execute
6	110	read and write
7	111	read, write, and execute

Access Permissions & UID/GID

- Operations vs. Permissions
 - Directory
 - X - pass through the dir (search bit), e.g., /usr/dict/words and PATH env var.
 - R - list of files under the dir.
 - W - update the dir, e.g., delete or create a file.
 - Everyone has the write permission for /tmp. Does it mean that you can delete any file in /tmp?
 - File
 - X - execute a file (which must be a regular file)
 - R - O_RDONLY or O_RDWR
 - W - O_WRONLY, O_RDWR, or O_TRUNC

Function – chmod & fchmod

```
#include <sys/types.h>
#include <sys/stat.h>
int chmod(const char *pathname, mode_t mode);
int fchmod(int filedes, mode_t mode);
int fchmodat(int fd, const char *pathname, mode_t mode,
             int flag);
```

- fchmod() is not in POSIX.1, but in SVR4/4.3+BSD
- Callers must be a superuser or effective UID = file UID.
- Mode = bitwise-OR (Fig 4.11).
 - S_I[RWX]USR,
 - S_I[RWS]GRP,
 - S_I[RWX]OTH,
 - S_ISVTX: sticky bit,
 - S_IS[UG]ID:set-usr-id/set-group-id.

Special Bits

- Sticky bit: 'T' or 't' on directory represents a sticky bit to protect files in the directory from misuse.
- Set-User-ID: 's' on executable files lead to user ID changes to provide additional privilege.

Sticky Bit

Permission on shared folders

- On systems, there are directories which are shared by many users.
- The users can use such directories to exchange files/information among processes owned different users.
- However,
 - with write permission to the directory for all the users in the group, one may delete files owned by other users.

Sticky Bit

- Sticky Bit (S_ISVTX) – saved-text bit
 - Not POSIX.1 – by SVR4 & 4.3+BSD
 - Only superusers can set it!
 - S_ISVTX directory file, e.g., /tmp
 - Remove/rename its file only if 'w' permission of the dir is set, and the process is belonging to superusers/owner of the file/dir.
 - S_ISVTX executable file
 - Used to save a copy of a S_ISVTX executable in the swap area to speed up the execution next time.
 - Not needed for a system with a virtual memory system and fast file system.

Set-Uid/set-Group-ID

The privilege of a process

```
int main()
{
    FILE *fp = fopen("./write.txt", "w+");
    testing_setbuf( fp );
    fclose( fp );
    return 0;
}
```

- The privilege of a process can be defined by the ID of the user starting the program.
- Single identity shows your royalty but limits the flexibility.



Real User ID vs. Effective User ID


- Effective User ID is mostly used for determining the permission for accessing a file.
- **Real User ID:**
 - When a user executes a program file, the real user ID is the user ID of the user.
- **Effective User ID:**
 - When a user executes a program file, the effective user ID is USUALLY the real user ID.
 - However, when SUID bit is set, the effective user ID is the owner of the program file.

/usr/bin/passwd

- /usr/bin/passwd:
 - user perms (root) - read, setuid on execute
 - group perms (sys) - read and execute/read and setgid on execute
 - others/anyone perms - read and execute

```
iMac:~ cshih$ ls -l /usr/bin/passwd
-r-sr-xr-x 1 root wheel 111968 8 18 2010 /usr/bin/passwd
iMac:~ cshih$
```

ruid/euid=
cshih/cshih



Set-User-ID/Set-Group-ID

- st_uid/st_gid: user/group ID of the owner of a file.
- How can a user change his/her password without having the write permission to /etc/passwd and /etc/shadow?

```
-rw-r--r--      1 root root 1746  2/21 13:00 /etc/passwd
-r-----      1 root root 1142  2/21 13:01 /etc/shadow
```

- A process could be associated with more than one ID.
 - Real user/group ID
 - Effective user/group ID
 - Supplementary group IDs
 - Saved set-user/group-ID

Example for chmod & fchmod

- Figure 4.12 – Page 107
 - chmod – updates on i-nodes
- For security reason, chmod() will clear set-group-ID
 - If the GID of a newly created file is not equal to the effective GID of the calling process (or one of the supplementary GID's), or the process is not a superuser, clear the set-group-ID bit!
 - if a non-superuser process writes to a set-uid/gid file.

Access Permissions & UID/GID

- File Access Test – each time a process creates/opens/deletes a file
 - If the effective UID == 0 → superuser!
 - If the effective UID == UID of the file
 - Check appropriate access permissions!
 - If the effective GID or any of its supplementary group* ID == GID of the file
 - Check appropriate access permissions!
 - Check appropriate access permissions for others!
- Related Commands: chmod & umask

* According IEEE 1003.1 – 1990 standard.

Ownership of new files

Ownership of a New File

- Rules:

- Owner UID of a file = the effective UID of the creating process
- Owner GID of a file – options under POSIX
 - GID of the file = the effective GID of the process
 - GID of the file = the GID of the residing dir
 - 4.3BSD and FIPS 151-1 always do it.
 - SVR4 needs to set the set-group-ID bit of the residing dir (mkdir)!

Linux

User's group as the group owner

```
csih@linux10:/tmp$ ls -ld /tmp
drwxrwxrwt 38 root root 765952 2011-04-13 05:43 /tmp/
csih@linux10:/tmp$ touch /tmp/new-file; ls -l /tmp/new-file
-rw-r--r-- 1 csih users 0 2011-04-13 05:43 /tmp/new-file
csih@linux10:/tmp$ ls -ld ~
drwxr-xr-x 23 csih users 1536 2011-04-13 05:36 /home/faculty/csih/
csih@linux10:/tmp$ touch ~/new-file; ls -l ~/new-file
-rw-r--r-- 1 csih users 0 2011-04-13 05:43 /home/faculty/csih/new-file
csih@linux10:/tmp$
```

BSD/OSX

Directory's group owner is that of new files

```
iMac:tmp csih$ ls -ld /tmp
lrwxr-xr-x@ 1 root wheel 11 11 10 2009 /tmp -> private/tmp
iMac:tmp csih$ touch /tmp/new-file; ls -l /tmp/new-file
-rw-r--r-- 1 csih wheel 0 4 13 05:44 /tmp/new-file
iMac:tmp csih$ ls -ld ~
drwxr-xr-x@ 139 csih csih 4726 4 13 05:38 /Users/csih
iMac:tmp csih$ touch ~/new-file; ls -l ~/new-file
-rw-r--r-- 1 csih csih 0 4 13 05:45 /Users/csih/new-file
iMac:tmp csih$
```

Function – umask

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
mode_t umask(mode_t cmask);
```

- The file mode to be turned OFF
 - cmask = bitwise-OR S_I[RWX]USR, etc (Figure 4.4).
 - Set permission to cmask & 0777 for directories
 - Set permission to cmask & 0666 for files
 - If cmask is 022, permission = 0666 & ~022 = 0644 = rw-r--r--
- The mask goes with the process only.
 - Inheritance from the parent!
- Program 4.9 – Page 104
 - umask

Permission for new file/dir

```
iMac:file cshih$ umask
0022
iMac:file cshih$ touch new-file; ls -l new-file
-rw-r--r--  1 cshih  cshih  0  4 13 06:07 new-file
iMac:file cshih$ rm -f new-file
iMac:file cshih$ umask 002
iMac:file cshih$ touch new-file; ls -l new-file
-rw-rw-r--  1 cshih  cshih  0  4 13 06:08 new-file
iMac:file cshih$ mkdir new-dir
iMac:file cshih$ umask
0002
iMac:file cshih$ ls -ld new-dir
drwxrwxr-x  2 cshih  cshih  68  4 13 06:10 new-dir
iMac:file cshih$
```

Function – chown, fchown, fchownat, lchown

```
#include <sys/types.h>
#include <unistd.h>
int chown(const char *pathname, uid_t owner, gid_t, grp);
int fchown(int filedes, uid_t owner, gid_t, grp);
int fchownat(int fd, const char *pathname, uid_t owner, gid_t
group,
int flag);
int lchown(const char *pathname, uid_t owner, gid_t, grp);
```

- lchown() is unique to SVR4. Under non-SVR4 systems, if the pathname to chown() is a symbolic link, only the ownership of the symbolic link is changed.
- -1 for owner or grp if no change is wanted.

Function – chown, fchown, fchownat, lchown

- _POSIX_CHOWN_RESTRICTED is in effect (check pathconf())
 - Superuser → the UID of the file can be changed!
 - The GID of the file can be changed if
 - the process owns the file, and
 - Parameter owner = UID of the file &
Parameter grp = the process GID or is in supplementary GID's
- set-user/group-ID bits would be cleared if chown is called by non-super users.

File Size

- File Sizes - `st_size`
 - Regular files - 0~max (`off_t`)
 - Directory files - multiples of 512 on Linux, or 34 on OSX.
 - Symbolic links - pathname length
 - /* a pipe's file size for SVR4 */
- File Holes
 - `st_blocks` vs `st_size` (`st_blksize`)
 - Many UNIX systems including OSX use 512bytes as the unit, not 1024 bytes.
 - Commands:
 - `"ls -l file.hole" == "wc -c file.hole"`
 - `du -s file.hole` → actual used size on disks, in numbers of blocks.
 - `cat -v file.hole > file.hole.copy`